AMD **Accelerated**
Parallel Processing
TECHNOLOGY

SAMPLE

Array vs ArrayView

1 Overview

1.1 Location $<APPSDKSamplesInstallPath>\samples\C++Amp\`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The default executables are placed in `$<APPSDKSamplesInstallPath>\samples\C++Amp\bin\x86\` for 32-bit builds, and `$<APPSDKSamplesInstallPath>\samples\C++Amp\bin\x86_64\` for 64-bit builds.

Type the following command(s).

1. `ArrayVsArrayView`

   This runs the program with the default options: s = (1024 * 768 * 3).

2. `ArrayVsArrayView -h`

   This prints the help file.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1   Command Line Options

Short Form	Long Form	Description
-h	--help	Show all command options and their respective meaning.
-q	--quiet	Quiet mode. Suppresses text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing.
-d	--deviceId	Select deviceId to be used (0 to N-1, where N is the ID of the device to be used).
-v	--version	AMD APP SDK version string.
-x	--samples	Number of example input values (multiples of three).

2 Introduction

When choosing between array and array_view in C++ AMP, arrays can be useful:

- Because DX interop APIs expect arrays as parameters.
- As staging buffers to optimize frequent data transfers that take place between the CPU and the GPU.
- To measure the performance of the data transfer to an accelerator.

Array vs ArrayView

- To copy to an accelerator asynchronously.

Array_views are useful for future-proofing code.

In this example, we explore staging buffers and how data transfers can be optimized using AMP arrays with them.

3 Implementation Details

As the current C++ AMP is implemented on top of DirectX 11, a copy between host buffer and device buffer consists of the following stages:

1. Create a CPU side staging buffer.
2. Copy from the host buffer to the mapped staging buffer.
3. Unmap the staging buffer, and perform a copy from the staging buffer to the GPU buffer.
4. Release the staging buffer.

Staging arrays are differentiated from normal arrays by their construction. A normal array is constructed by specifying the accelerator_view on which the array is to be created. If this parameter is not provided, the array is created on the default accelerator.

A staging array is constructed with a second accelerator_view acclView2 (along with the accelerator_view acclView1, which specifies the device on which the array is to be created), where:

- The stagingArray is physically located on acclView1. So the "accelerator_view" property of stagingArray returns the value of acclView1.
- The second accelerator_view parameter is called the associated accelerator_view. It is a hint to the runtime that a high frequency of data transfers between this array and other arrays on acclView2 will occur.

The C++ AMP runtime uses this hint so when the stagingArray is to be copied to another array on deviceAcclView, it avoids an additional copy, buffer creation, and releases; thus saving valuable time. Currently, it only helps in the case where acclView1 is a CPU accelerator_view, and acclView2 is a GPU accelerator_view.

This example shows how C++ AMP staging arrays can be created and used efficiently with a noticeable performance boost over array_view. In this example, a simple color conversion routine of RGB to YUV 4:4:4 is used to bring out the uses and advantages of staging arrays.

4 Recommended Input Option Settings

For best performance, enter the following on the command line: `-x 2359296 -q -t -e`

5 References

1. http://www.danielmoth.com/Blog/array-And-Arrayview-From-Amph.aspx
2. http://blogs.msdn.com/b/nativeconcurrency/archive/2012/07/17/choosing-between-array-and-array-view-in-c-amp.aspx

3. http://blogs.msdn.com/b/nativeconcurrency/archive/2011/11/10/staging-arrays-in-c-amp.aspx

4. http://en.wikipedia.org/wiki/YUV

AMD